

DevOps Exercise

The files checked into my GitHub repository: <https://github.com/yellowzoneallan/buildit>

Dockerfile	Buildit code as docker image
README.md	Initial commit
Vagrantfile	Vagrant config
docker-compose.yml	Docker Compose Config
nginx.lb.conf	Nginx conf to be a load balancer 80 -> dockers:3000
setup.sh	Build the buildit system

The docker images in my solution were:

- git - a private GitHub for testing
- Jenkins - a private Jenkins for testing
- builditapp - the example code wrapped in an image
- nginx - nginx importing a config file to covert it to a load balancer
- rusmckendrick/ab - random prepackaged apache benchmark image

This alias was setup so "localgit" commands were identical to "git" commands

```
alias localgit='sudo docker run -ti --rm -v ${HOME}:/root -v $(pwd):/git alpine/git "$@"'
```

The Vagrant and docker-compose file combination allows browser access to 192.168.3.3

Docker-compose allows the number of instances to vary, in this example 10 were started and apache benchmark shows 10,000 successful hits in 5.6s.

```
vagrant@ubuntu-eoan:~$ sudo docker-compose up --scale builditapp=10 -d
```

```
Creating network "vagrant_default" with the default driver
```

```
Creating vagrant_builditapp_1 ... done
```

```
Creating vagrant_builditapp_2 ... done
```

```
Creating vagrant_builditapp_3 ... done
```

```
Creating vagrant_builditapp_4 ... done
```

```
Creating vagrant_builditapp_5 ... done
```

```
Creating vagrant_builditapp_6 ... done
```

```
Creating vagrant_builditapp_7 ... done
```

```
Creating vagrant_builditapp_8 ... done
```

```
Creating vagrant_builditapp_9 ... done
```

```
Creating vagrant_builditapp_10 ... done
```

```
Creating vagrant_jenkins_1 ... done
```

```
Creating vagrant_localgit_1 ... done
```

```
Creating vagrant_nginx_1 ... done
```

```
vagrant@ubuntu-eoan:~$ sudo docker run rusmckendrick/ab ab -k -n 10000 -c 16 http://192.168.3.3/
```

```
This is ApacheBench, Version 2.3 <$Revision: 1826891 $>
```

```
Finished 10000 requests
```

```
Server Software:      nginx/1.17.9
```

```
Server Hostname:      192.168.3.3
```

```
Server Port:          80
```

```
Document Path:        /
```

```
Document Length:      44 bytes
```

```
Concurrency Level:    16
```

```
Time taken for tests:  5.672 seconds
```

Complete requests: 10000
Failed requests: 0
Keep-Alive requests: 0
Total transferred: 1410000 bytes
HTML transferred: 440000 bytes
Requests per second: 1763.04 [#/sec] (mean)
Time per request: 9.075 [ms] (mean)
Time per request: 0.567 [ms] (mean, across all concurrent requests)
Transfer rate: 242.76 [Kbytes/sec] received

Connection Times (ms)

	min	mean[+/-sd]	median	max
Connect:	0	1 0.9	0	7
Processing:	1	8 4.0	8	66
Waiting:	0	8 4.1	8	66
Total:	1	9 3.8	8	66

WARNING: The median and mean for the initial connection time are not within a normal deviation
These results are probably not that reliable.

Percentage of the requests served within a certain time (ms)

50%	8
90%	11
99%	24
100%	66 (longest request)

vagrant@ubuntu-eoan:~\$

Platform Engineer task

There already exists an excellent open source product which achieves the task objectives and much more. The application is packaged as a java jar, which can be scaled effectively by selecting larger VMs, JVMs and the java threadpooling, garbage collection and heap tuning parameters.

Via a browser on your laptop <https://darcyripper.com/features/downloads/>
Then extract the zip and click the darcyripper.jar, it is obvious what to do.

If the exercise is about automation here's the ubuntu linux commands needed:

```
sudo curl --insecure -L "https://darcyripper.com/download/1230/darcyripper.zip" -o  
darcyripper.zip # download ripper  
sudo apt install -y openjdk-14-jre-headless # install java  
sudo apt install -y unzip # install unzip  
unzip darcyripper.zip; cd darcyripper/ # unzip ripper
```